

HW 2

By Zhuo Leng

step 1: read data

```
In [194]: #import read.py  
import read
```

```
In [36]: df = read.read_data('credit-data.csv')  
df.head()
```

```
Out[36]:
```

	PersonID	SeriousDlqin2yrs	RevolvingUtilizationOfUnsecuredLines	age	zipcode	Num 59D
0	1	1	0.766127	45	60644	2
1	2	0	0.957151	40	60637	0
2	3	0	0.658180	38	60601	1
3	4	0	0.233810	30	60601	0
4	5	0	0.907239	49	60625	1

step 2: explore data

```
In [205]: #import explore.py  
%matplotlib inline  
import explore
```

```
In [16]: df.describe()
```

```
Out[16]:
```

	PersonID	SeriousDlqin2yrs	RevolvingUtilizationOfUnsecuredLines	age
count	150000.000000	150000.000000	150000.000000	150000.000000
mean	75000.500000	0.066840	6.048438	52.295200
std	43301.414527	0.249746	249.755371	14.771860
min	1.000000	0.000000	0.000000	0.000000
25%	37500.750000	0.000000	0.029867	41.000000
50%	75000.500000	0.000000	0.154181	52.000000
75%	112500.250000	0.000000	0.559046	63.000000
max	150000.000000	1.000000	50708.000000	109.000000


Type Markdown and LaTeX: α^2

```
In [17]: explore.number_count(df, 'NumberOfDependents')
```

```
Out[17]: 0.0    86902
          1.0    26316
          2.0    19522
          3.0     9483
          4.0     2862
          5.0      746
          6.0     158
          7.0      51
          8.0      24
          9.0       5
Name: NumberOfDependents, dtype: int64
```

```
In [18]: pd.value_counts(df.NumberOfDependents, ascending=False).head(10).plot(kind
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x1127b9b38>
```




From the top 10 number account of NumberOfDependents, we could know more than 1/2 of family has 0.0 number of dependents in family excluding themselves(spouse, children etc.). They live on their own.Number of dependents in family excluding themselves (spouse, children etc.)

```
In [19]: explore.number_count(df, 'age')
```

```
Out[19]: 49      3837
         48      3806
         50      3753
         63      3719
         47      3719
         46      3714
         53      3648
         51      3627
         52      3609
         56      3589
         Name: age, dtype: int64
```

```
In [20]: pd.value_counts(df.age, ascending=False).head(10).plot(kind = 'bar')
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x109849b38>
```




From the top 10 number account of age, we could know the age of people don't different a lot. Most people are at the age of 49.

```
In [21]: explore.number_count(df, 'zipcode')
```

```
Out[21]: 60625      33514
        60629      16840
        60637      16625
        60657      16624
        60618      16612
        60804      16605
        60601      16599
        60644      16581
        Name: zipcode, dtype: int64
```

```
In [22]: pd.value_counts(df.zipcode, ascending=False).plot(kind = 'bar')
```

```
Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x111f7a160>
```



Take a look at zipcode variable, most people live in zipcode area 60625.

cross tab


```
In [23]: explore.crosstable(df, 'NumberOfTimes90DaysLate', 'SeriousDlqin2yrs')
```

```
Out[23]:
```

SeriousDlqin2yrs	0	1
NumberOfTimes90DaysLate		
0	135108	6554
1	3478	1765
2	779	776
3	282	385
4	96	195
5	48	83
6	32	48
7	7	31
8	6	15
9	5	14
10	3	5
11	2	3
12	1	1
13	2	2
14	1	1
15	2	0
17	0	1
96	1	4
98	121	143

```
In [203]: pd.crosstab(df.NumberOfTimes90DaysLate, df.SeriousDlqin2yrs).plot(figsize
```

```
Out[203]: <matplotlib.axes._subplots.AxesSubplot at 0x127c9fc88>
```



From the cross table of NumberOfTimes90DaysLate and SeriousDlqin2yrs, most people experienced 90 days past due delinquency or worse have more number of times borrower has been 90 days or more past due. People do not experienced 90 days past due delinquency or worse have less number of times borrower has been 90 days or more past due.

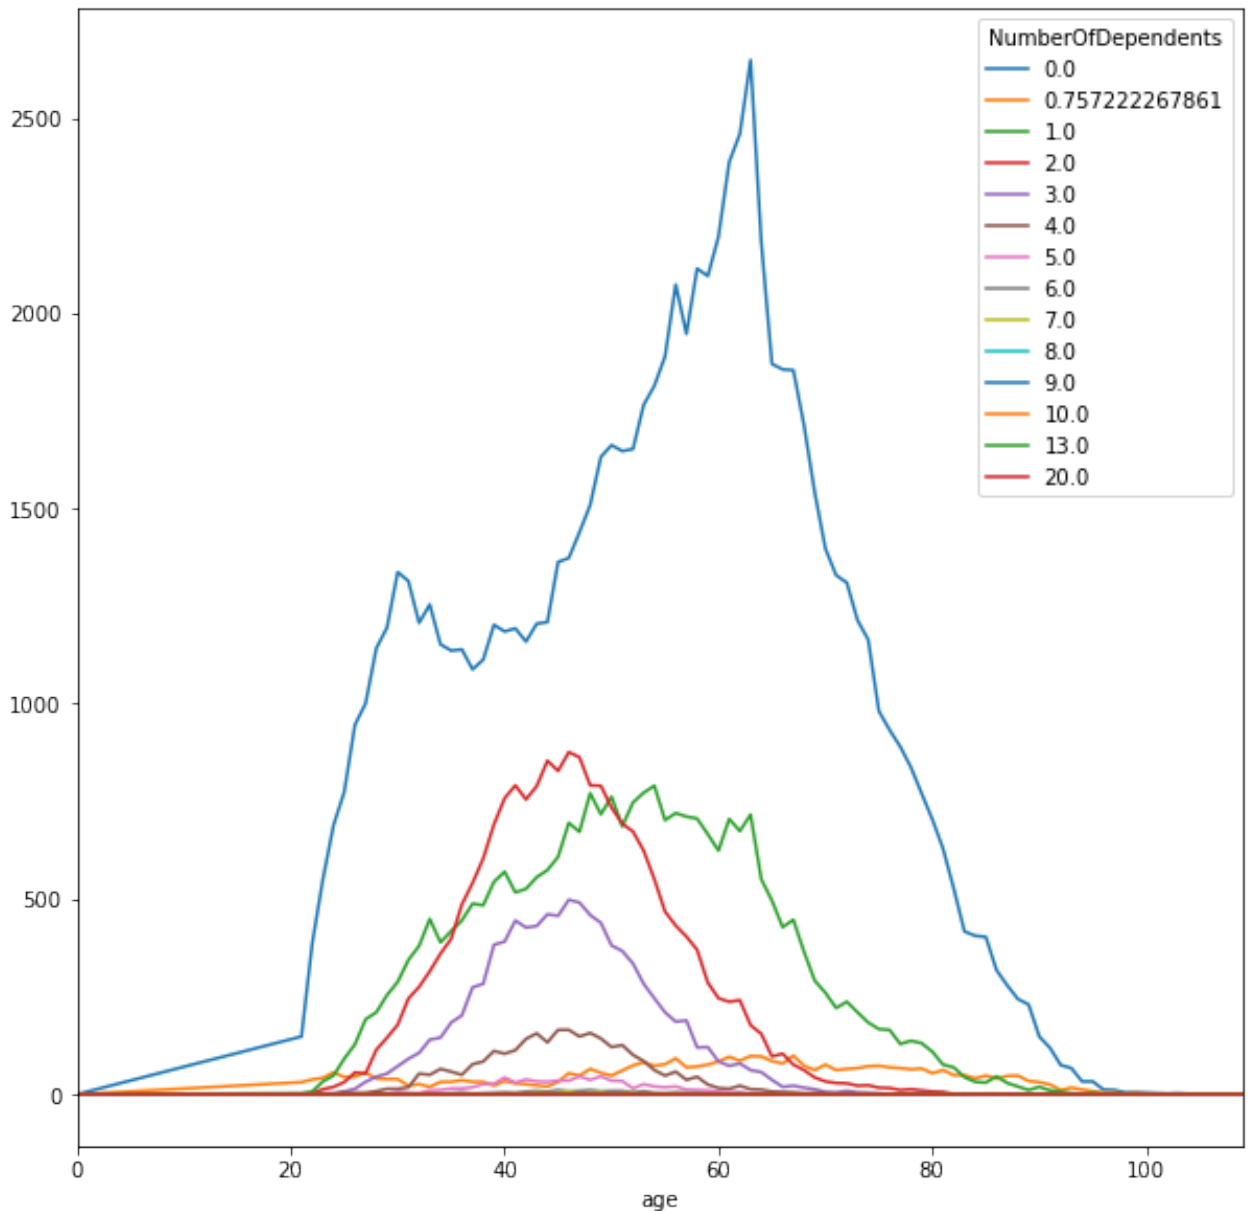
```
In [25]: pd.crosstab(df.age, df.NumberOfDependents)
```

```
Out[25]:
```

NumberOfDependents	0.0	1.0	2.0	3.0	4.0	5.0	6.0	7.0	8.0	9.0	10.0	13.0	20.0
age													
0	0	0	1	0	0	0	0	0	0	0	0	0	0
21	148	3	1	0	0	0	0	0	0	0	0	0	0
22	385	7	2	2	0	0	0	0	0	0	0	0	0
23	550	33	13	3	0	0	0	0	0	0	0	0	0
24	689	48	19	3	1	0	0	0	0	0	0	0	0
25	774	91	31	7	5	1	0	0	0	0	0	0	0
26	946	128	56	14	4	0	0	0	0	0	0	0	0
27	1001	192	53	32	4	0	0	0	0	0	0	0	0
28	1142	210	114	45	8	1	0	0	0	0	0	0	0


```
In [202]: pd.crosstab(df.age, df.NumberOfDependents).plot(figsize=(10,10))
```

```
Out[202]: <matplotlib.axes._subplots.AxesSubplot at 0x1207f2940>
```



From the cross table of age and NumberOfDependents, most people have 9 number of dependents in family excluding themselves (spouse, children etc.). Also, we could know, with the increase of ages, number of dependents in family excluding themselves (spouse, children etc.) will increase. However, when after age 40-60, number of dependents in family excluding themselves (spouse, children etc.) will decrease.

step 3: pre-processing data: handing missing value

```
In [ ]: #import PreProcess.py
import PreProcess
```

```
In [37]: df_lng = pd.melt(df)
df_lng.head()
```

```
Out[37]:
```

	variable	value
0	PersonID	1.0
1	PersonID	2.0
2	PersonID	3.0
3	PersonID	4.0
4	PersonID	5.0

```
In [38]: PreProcess.print_null_freq(df)
```

```
Out[38]:
```

value	False	True
variable		
DebtRatio	150000	0
MonthlyIncome	120269	29731
NumberOfDependents	146076	3924
NumberOfOpenCreditLinesAndLoans	150000	0
NumberOfTime30-59DaysPastDueNotWorse	150000	0
NumberOfTime60-89DaysPastDueNotWorse	150000	0
NumberOfTimes90DaysLate	150000	0
NumberRealEstateLoansOrLines	150000	0
PersonID	150000	0
RevolvingUtilizationOfUnsecuredLines	150000	0
SeriousDlqin2yrs	150000	0
age	150000	0
zipcode	150000	0

From the crosstable table, we could know that MonthlyIncome and NumberOfDependents are two variables with missing value. We need to fill them with mean value.

```
In [39]: df.MonthlyIncome.describe()
```

```
Out[39]: count      1.202690e+05  
mean        6.670221e+03  
std         1.438467e+04  
min         0.000000e+00  
25%         3.400000e+03  
50%         5.400000e+03  
75%         8.249000e+03  
max         3.008750e+06  
Name: MonthlyIncome, dtype: float64
```

```
In [40]: #fill MonthlyIncome na with mean  
PreProcess.fill_na(mean,df,'MonthlyIncome')
```

```
In [41]: df.NumberOfDependents.describe()
```

```
Out[41]: count      146076.000000  
mean           0.757222  
std            1.115086  
min            0.000000  
25%            0.000000  
50%            0.000000  
75%            1.000000  
max            20.000000  
Name: NumberOfDependents, dtype: float64
```

```
In [42]: #fill MonthlyIncome na with mean  
PreProcess.fill_na(mean,df,'NumberOfDependents')
```

```
In [48]: #check again for null value by using cross table  
PreProcess.print_null_freq(df)
```

Out[48]:

value	False
variable	
DebtRatio	150000
MonthlyIncome	150000
NumberOfDependents	150000
NumberOfOpenCreditLinesAndLoans	150000
NumberOfTime30-59DaysPastDueNotWorse	150000
NumberOfTime60-89DaysPastDueNotWorse	150000
NumberOfTimes90DaysLate	150000
NumberRealEstateLoansOrLines	150000
PersonID	150000
RevolvingUtilizationOfUnsecuredLines	150000
SeriousDlqin2yrs	150000
age	150000
zipcode	150000

Now all the variables have 150000 number of no-null value. So we successfully fill the missing value with mean of relevant variables. After clean the data, we need to yield a new csv.

```
In [49]: df.to_csv("credit-data-post-import.csv", index=False)
```

step 4: Generate Features/Predictors

```
In [64]: #import feature.py
import feature
df.head()
```

```
Out[64]:
```

loans	NumberOfTimes90DaysLate	NumberRealEstateLoansOrLines	NumberOfTime60-89DaysPastDueNotW
	0	6	0
	0	0	0
	1	0	0
	0	0	0
	0	1	0

```
In [70]: #categorical features
#SeriousDlqin2yrs
df['SeriousDlqin2yrs_False'] = feature.categorical_var(df, 'SeriousDlqin2y
df['SeriousDlqin2yrs_True'] = feature.categorical_var(df, 'SeriousDlqin2yr
feature.categorical_var(df, 'SeriousDlqin2yrs')
```

```
Out[70]:
```

	0	1
0	0	1
1	1	0
2	1	0
3	1	0
4	1	0
5	1	0
6	1	0
7	1	0
8	1	0
9	1	0

```
In [78]: #discretize continuous variables
#RevolvingUtilizationOfUnsecuredLines
max(df['RevolvingUtilizationOfUnsecuredLines']) - min(df['RevolvingUtiliza
```

```
Out[78]: 50708.0
```

```
In [83]: UtilizationOfUnsecuredLines) = 50708, so we choose bin = 100
         edLines_dis'] = feature.discretize_continuous_var(df, 'RevolvingUtilization
         edLines_dis']
149984      (-50.708, 507.08]
149985      (-50.708, 507.08]
149986      (-50.708, 507.08]
149987      (-50.708, 507.08]
149988      (-50.708, 507.08]
149989      (-50.708, 507.08]
149990      (-50.708, 507.08]
149991      (-50.708, 507.08]
149992      (-50.708, 507.08]
149993      (-50.708, 507.08]
149994      (-50.708, 507.08]
149995      (-50.708, 507.08]
149996      (-50.708, 507.08]
149997      (-50.708, 507.08]
149998      (-50.708, 507.08]
149999      (-50.708, 507.08]
Name: RevolvingUtilizationOfUnsecuredLines_dis, dtype: category
Categories (100, object): [(-50.708, 507.08] < (507.08, 1014.16] < (10
14.16, 1521.24] < (1521.24, 2028.32] ... (48679.68, 49186.76] < (49186
.76, 49693.84] < (49693.84, 50200.92] < (50200.92, 50708]]
```

```
In [75]: #discretize continuous variables
         #age
         max(df['age']) - min(df['age'])
```

```
Out[75]: 109
```

```
In [76]: #because age varibale (max-min) =109, we choose bins = 20
df['age_dis'] = feature.discretize_continuous_var(df,'age', 20)
df['age_dis']
```

```
Out[76]: 0      (43.6, 49.05]
1      (38.15, 43.6]
2      (32.7, 38.15]
3      (27.25, 32.7]
4      (43.6, 49.05]
5      (70.85, 76.3]
6      (54.5, 59.95]
7      (38.15, 43.6]
8      (21.8, 27.25]
9      (54.5, 59.95]
10     (27.25, 32.7]
11     (49.05, 54.5]
12     (43.6, 49.05]
13     (38.15, 43.6]
14     (70.85, 76.3]
15     (59.95, 65.4]
16     (76.3, 81.75]
17     (49.05, 54.5]
18     (38.15, 43.6]
19     (21.8, 27.25]
```

```
In [92]: #categorical features
#zipcode
zip = feature.categorical_var(df,'zipcode')
zip
```

```
Out[92]:
```

	60601	60618	60625	60629	60637	60644	60657	60804
0	0	0	0	0	0	1	0	0
1	0	0	0	0	1	0	0	0
2	1	0	0	0	0	0	0	0
3	1	0	0	0	0	0	0	0
4	0	0	1	0	0	0	0	0
5	0	0	0	1	0	0	0	0
6	0	0	0	0	1	0	0	0
7	0	0	1	0	0	0	0	0
8	0	0	0	0	0	0	0	1
9	0	0	0	1	0	0	0	0

```
In [95]: df['zip_60601'] = zip.iloc[:,0]
df['zip_60618'] = zip.iloc[:,1]
df['zip_60625'] = zip.iloc[:,2]
df['zip_60629'] = zip.iloc[:,3]
df['zip_60637'] = zip.iloc[:,4]
df['zip_60644'] = zip.iloc[:,5]
df['zip_60657'] = zip.iloc[:,6]
df['zip_60804'] = zip.iloc[:,7]
```

```
In [105]: #categorical features
#NumberOfTime30-59DaysPastDueNotWorse
num = pd.get_dummies(df.iloc[:,5], prefix = 'NumberOfTime30-59DaysPastDue
num
```

Out[105]:

	NumberOfTime30-59DaysPastDueNotWorse_0	NumberOfTime30-59DaysPastDueNotWorse_1	NumberOfTime30-59DaysPastDueNotWors
0	0	0	1
1	1	0	0
2	0	1	0
3	1	0	0
4	0	1	0
5	1	0	0
6	1	0	0
7	1	0	0
8	1	0	0
9	1	0	0


```
In [114]: #concat variables to df
df = pd.concat([df, num], axis=1)
df
```

Out[114]:

	PersonID	SeriousDlqin2yrs	RevolvingUtilizationOfUnsecuredLines	age	zipcode
0	1	1	0.766127	45	60644
1	2	0	0.957151	40	60637
2	3	0	0.658180	38	60601
3	4	0	0.233810	30	60601
4	5	0	0.907239	49	60625
5	6	0	0.213179	74	60629
6	7	0	0.305682	57	60637
7	8	0	0.754464	39	60625
8	9	0	0.116951	27	60804
9	10	0	0.189169	57	60629

```
In [115]: #discretize continuous variables
#DebtRatio
max(df['DebtRatio']) - min(df['DebtRatio'])
```

Out[115]: 329664.0

```
In [116]: #because age varibale (max-min) =329664, we choose bins = 100  
df['DebtRatio_dis'] = feature.discretize_continuous_var(df,'age', 100)  
df['age_dis']
```

```
Out[116]: 0      (43.6, 49.05]  
1      (38.15, 43.6]  
2      (32.7, 38.15]  
3      (27.25, 32.7]  
4      (43.6, 49.05]  
5      (70.85, 76.3]  
6      (54.5, 59.95]  
7      (38.15, 43.6]  
8      (21.8, 27.25]  
9      (54.5, 59.95]  
10     (27.25, 32.7]  
11     (49.05, 54.5]  
12     (43.6, 49.05]  
13     (38.15, 43.6]  
14     (70.85, 76.3]  
15     (59.95, 65.4]  
16     (76.3, 81.75]  
17     (49.05, 54.5]  
18     (38.15, 43.6]  
19     (21.8, 27.25]
```

```
In [117]: #discretize continuous variables  
#MonthlyIncome  
max(df['MonthlyIncome']) - min(df['MonthlyIncome'])
```

```
Out[117]: 3008750.0
```

```
In [118]: #because age varibale (max-min) =329664, we choose bins = 1000
df['MonthlyIncome_dis'] = feature.discretize_continuous_var(df, 'MonthlyIncome')
df['MonthlyIncome_dis']
```

```
Out[118]: 0          (9026.25, 12035]
1      (-3008.75, 3008.75]
2          (3008.75, 6017.5]
3          (3008.75, 6017.5]
4      (63183.75, 66192.5]
5          (3008.75, 6017.5]
6          (6017.5, 9026.25]
7          (3008.75, 6017.5]
8          (6017.5, 9026.25]
9      (21061.25, 24070]
10     (-3008.75, 3008.75]
11     (6017.5, 9026.25]
12     (12035, 15043.75]
13     (12035, 15043.75]
14     (-3008.75, 3008.75]
15     (9026.25, 12035]
16     (6017.5, 9026.25]
17     (6017.5, 9026.25]
18     (3008.75, 6017.5]
19     (-3008.75, 3008.75]
```

```
In [127]: #categorical features
#NumberOfDependents
num = feature.categorical_var(df, 'NumberOfDependents')
num
```

Out[127]:

	0.0	0.757222267861	1.0	2.0	3.0	4.0	5.0	6.0	7.0	8.0	9.0	10.0	13.0	20.0
0	0	0	0	1	0	0	0	0	0	0	0	0	0	0
1	0	0	1	0	0	0	0	0	0	0	0	0	0	0
2	1	0	0	0	0	0	0	0	0	0	0	0	0	0
3	1	0	0	0	0	0	0	0	0	0	0	0	0	0
4	1	0	0	0	0	0	0	0	0	0	0	0	0	0
5	0	0	1	0	0	0	0	0	0	0	0	0	0	0
6	1	0	0	0	0	0	0	0	0	0	0	0	0	0
7	1	0	0	0	0	0	0	0	0	0	0	0	0	0
8	0	1	0	0	0	0	0	0	0	0	0	0	0	0
9	0	0	0	1	0	0	0	0	0	0	0	0	0	0

```
In [137]: num = num.add_prefix('NumberOfDependents')
num
```

149990	0	0
149991	1	0
149992	0	0
149993	1	0
149994	1	0
149995	1	0
149996	0	0
149997	1	0
149998	1	0
149999	1	0

150000 rows × 14 columns

```
In [138]: #concat variables to df
df = pd.concat([df, num], axis=1)
df
```

Out[138]:

	PersonID	SeriousDlqin2yrs	RevolvingUtilizationOfUnsecuredLines	age	zipcode
0	1	1	0.766127	45	60644
1	2	0	0.957151	40	60637
2	3	0	0.658180	38	60601
3	4	0	0.233810	30	60601
4	5	0	0.907239	49	60625
5	6	0	0.213179	74	60629
6	7	0	0.305682	57	60637
7	8	0	0.754464	39	60625
8	9	0	0.116951	27	60804
9	10	0	0.189169	57	60629

step 5: build classifier: Logistic Regression

The first model I try to build is a logistic regression model

```
In [164]: #import model.py
import model
```

```
In [149]: #the feature I choose for my first model is as below:
#I choose to use all variables in this model, except personID as my feature
# 'SeriousDlqin2yrs' is my dependent variables
#features
features = df.columns[2:13]
features
```

```
Out[149]: Index(['RevolvingUtilizationOfUnsecuredLines', 'age', 'zipcode',
                'NumberOfTime30-59DaysPastDueNotWorse', 'DebtRatio', 'MonthlyIncome',
                'NumberOfOpenCreditLinesAndLoans', 'NumberOfTimes90DaysLate',
                'NumberRealEstateLoansOrLines', 'NumberOfTime60-89DaysPastDueNotWorse',
                'NumberOfDependents'],
                dtype='object')
```

```
In [186]: df[features]
```

```
Out[186]:
```

	RevolvingUtilizationOfUnsecuredLines	age	zipcode	NumberOfTime30-59DaysPastDueNotWorse	DebtRatio
0	0.766127	45	60644	2	0.0681
1	0.957151	40	60637	0	0.0380
2	0.658180	38	60601	1	0.0391
3	0.233810	30	60601	0	0.0312
4	0.907239	49	60625	1	0.0391
5	0.213179	74	60629	0	0.0312
6	0.305682	57	60637	0	0.0571
7	0.754464	39	60625	0	0.0312
8	0.116951	27	60804	0	0.0461
9	0.189169	57	60629	0	0.0312

```
In [151]: outcome_var = df.columns[1]
outcome_var
```

```
Out[151]: 'SeriousDlqin2yrs'
```

```
In [160]: #build model
lm = model.logistic_regression(df, features, outcome_var)
lm
```

```
Out[160]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept
= True,
            intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs
= 1,
            penalty='l2', random_state=None, solver='liblinear', tol=0.0
001,
            verbose=0, warm_start=False)
```

```
In [171]: print(lm.coef_)

[[ 7.77870825e-06 -2.90907430e-02 -1.70008042e-05  1.22972174e-02
 -1.53991320e-05 -2.80009662e-05 -1.29996122e-03  1.08649202e-02
  4.97825084e-04  8.57494163e-03  1.69177367e-03]]
```

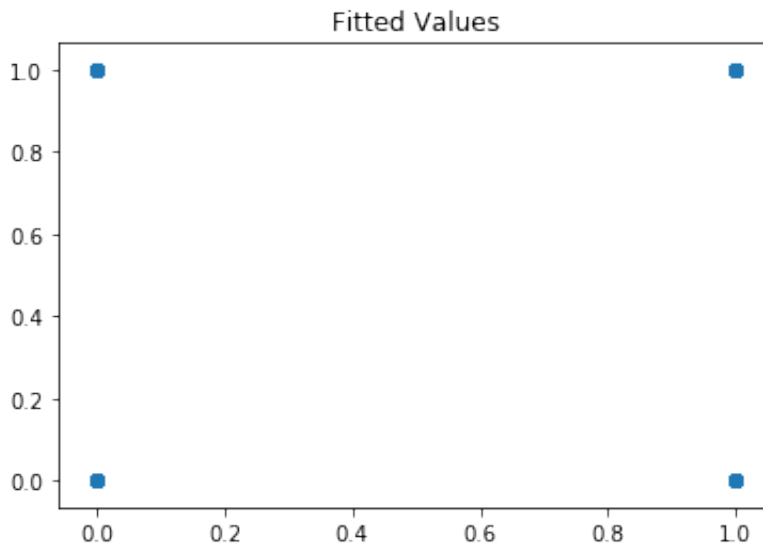
```
In [172]: # predict
lm.predict(df[features])
```

```
Out[172]: array([0, 0, 0, ..., 0, 0, 0])
```

```
In [180]: # make predictions
expected = df[outcome_var]
predicted = lm.predict(df[features])
```

```
In [182]: #to see the predicted value vs. actual value  
model.plot_pred_actual(df, features, outcome_var)
```

```
Out[182]: <matplotlib.text.Text at 0x125cb1eb8>
```



step 6: Evaluate Classifier

```
In [189]: import evaluation
```

```
In [185]: #Explained variance score: 1 is perfect prediction  
#and 0 means that there is no linear relationship  
#between X and y.  
evaluation.accuracy(lm, df, features, outcome_var)
```

```
Out[185]: 0.93332000000000004
```

From the accuracy score above, we could know there are strong linear relationship between feature and outcome.

```
In [192]: # summarize the fit of the model by using model.mse  
mse = model.mse(lm, df, features, outcome_var)  
mse
```

```
Out[192]: 0.066680000000000003
```

```
In [193]: print (lm.intercept_, lm.coef_, mse)

[ -1.87112635e-07] [[ 7.77870825e-06 -2.90907430e-02 -1.70008042e-0
5 1.22972174e-02
-1.53991320e-05 -2.80009662e-05 -1.29996122e-03 1.08649202e-02
4.97825084e-04 8.57494163e-03 1.69177367e-03]] 0.06668
```

The mse of this model is 0.066680000000000003, which is not very large. However, we still need compare it with other model's mse in future study.

```
In [ ]:
```